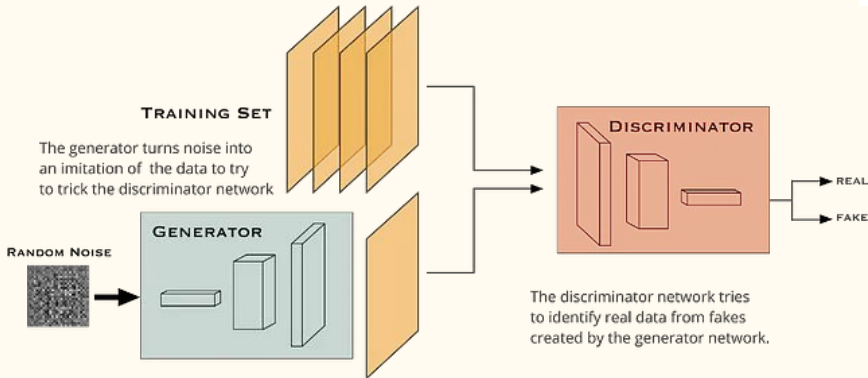
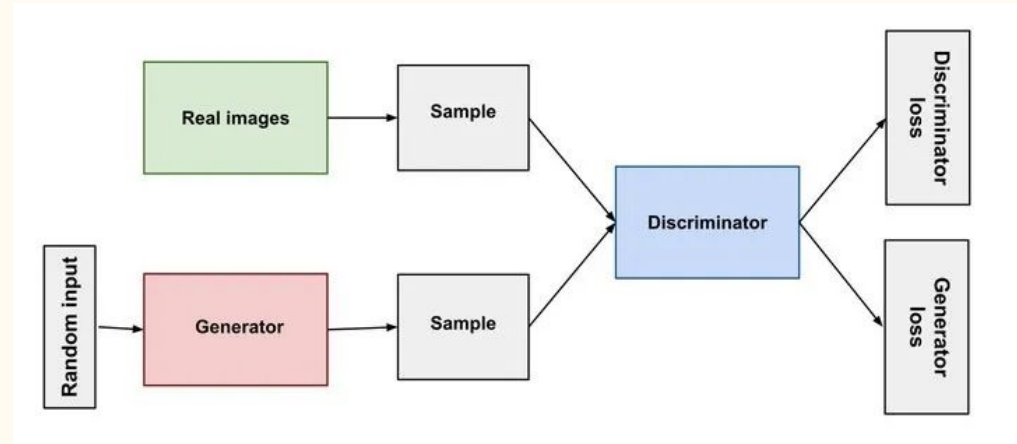


Single Image SRGAN

—

By: James Cheung


Generative Adversarial Networks




Generative Adversarial Networks Steps

1. From random Noise seed (latent vector) a Generator Network creates fake images
2. Real images are loaded
3. Images both real and fake are past through a Discriminator Network
4. Discriminator Network makes a binary classification decision trying to distinguish between real and fake images
5. Step 1-4 is looped and repeated: Re-trains the Generator Network / Discriminator Network depending on who wins and goes on for user-determined number of epochs
6. Generator model is saved to create new, realistic fake data

Generator Network



Discriminator Network



Applications of GAN:


- StackGAN - text-to-image synthesis (e.g. DALL-E)
- DCGAN - generate fake human faces
- Conditional GANs - generate images of people aging
- CycleGAN - Photoblending
- CycleGAN - Image Style Transfer
- Medical Image Synthesis using GANs for Pulmonary Chest X-rays



(a)



(b)




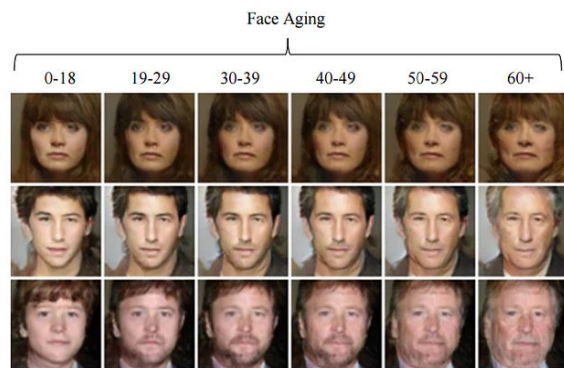
(c)



(d)



this bird is red with white and has a very short beak



Super Resolution Generative Adversarial Network

- Using CNN and VGG19 in the model
- Used 250 pictures and 40 epochs
- Image from 32x32 to upscale to 128x128 i.e. 4x magnification
- Instead of calculating loss with Mean-Squared-Error which results in overly smooth textures, peceptual loss function is used:
 - Content loss - uses Euclidean distance between feature representations
 - Adversarial loss - based on probability of the discriminator over all trainig samples

$$l^{SR} = \underbrace{l_X^{SR}}_{\text{content loss}} + \underbrace{10^{-3}l_{Gen}^{SR}}_{\text{adversarial loss}}$$

perceptual loss (for VGG based content losses)

Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network

Christian Ledig, Lucas Theis, Ferenc Huszar, Jose Caballero, Andrew Cunningham, Alejandro Acosta, Andrew Aitken, Atykhan Tejani, Johannes Totz, Zehan Wang, Wenzhe Shi
Twitter

{cledig,ltheis,fhuszar,jcaballero,aacostadiaz,aitken,atejani,jtotz,zehang,wshi}@twitter.com

Abstract

Despite the breakthroughs in accuracy and speed of single image super-resolution using faster and deeper convolutional neural networks, one central problem remains largely unsolved: how do we recover the finer texture details when we super-resolve to large upscaling factors? The behavior of optimization-based super-resolution methods is principally driven by the choice of the objective function. Recent work has largely focused on minimizing the mean squared reconstruction error. The resulting estimates have high peak signal-to-noise ratios, but they are often lacking high-frequency details and are perceptually unsatisfying in the sense that they fail to match the fidelity expected at the higher resolution. In this paper, we present SRGAN, a generative adversarial network (GAN) for image super-resolution (SR). To our knowledge, it is the first framework capable of inferring photo-realistic natural images for 4x upscaling factors. To achieve this, we propose a perceptual loss function which consists of an adversarial loss and a content loss. The adversarial loss pushes our solution to the natural image manifold using a discriminator network that is trained to differentiate between the super-resolved images and original photo-realistic images. In addition, we use a content loss motivated by perceptual similarity instead of similarity in pixel space. Our deep residual network is able to recover photo-realistic textures from heavily downsampled images on public benchmarks. An extensive mean-opinion-score (MOS) test shows hugely significant gains in perceptual quality using SRGAN. The MOS scores obtained with SRGAN are closer to those of the original high-resolution images than to those obtained with any state-of-the-art method.

1. Introduction

The highly challenging task of estimating a high-resolution (HR) image from its low-resolution (LR) counterpart is referred to as super-resolution (SR). SR received substantial attention from within the computer vision research community and has a wide range of applications [63, 71, 42].


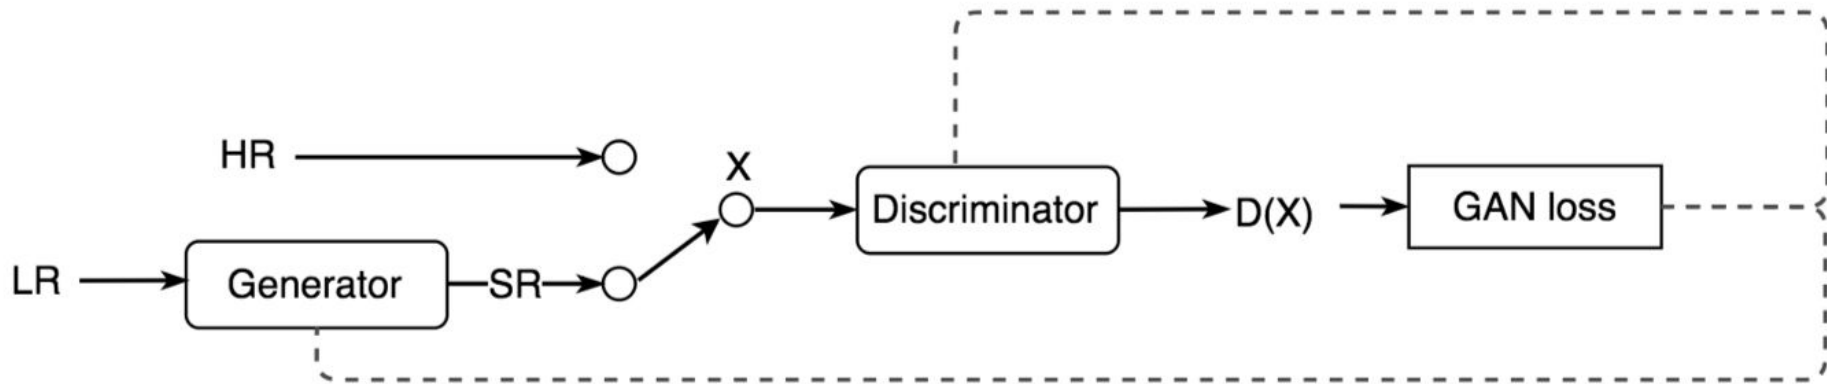


Figure 1: Super-resolved image (left) is almost indistinguishable from original (right). [4x upscaling]

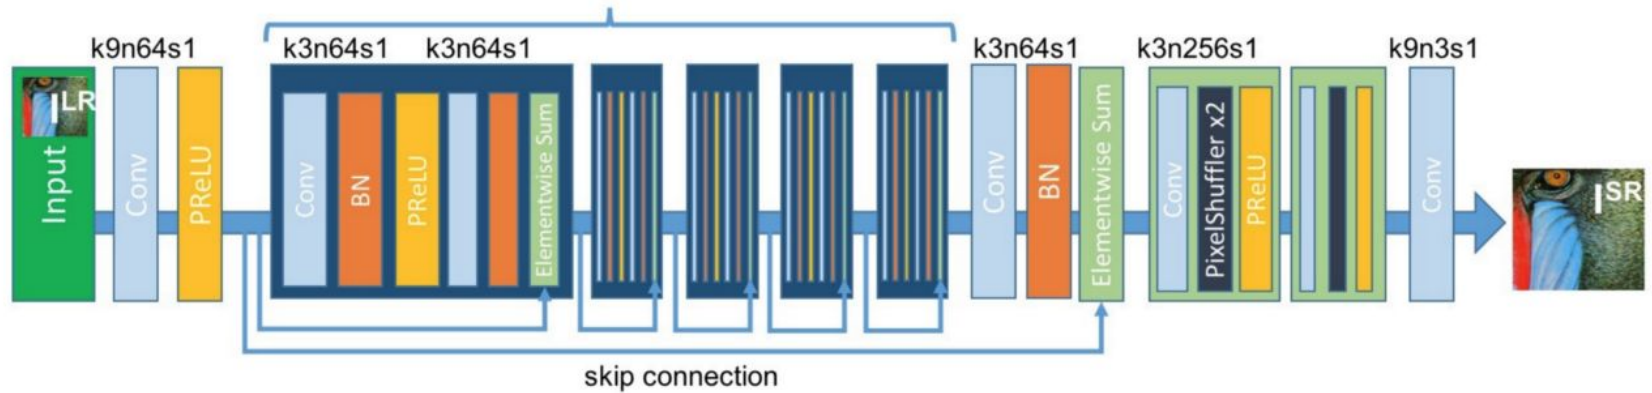
The ill-posed nature of the underdetermined SR problem is particularly pronounced for high upscaling factors, for which texture detail in the reconstructed SR images is typically absent. The optimization target of supervised SR algorithms is commonly the minimization of the mean squared error (MSE) between the recovered HR image and the ground truth. This is convenient as minimizing MSE also maximizes the peak signal-to-noise ratio (PSNR), which is a common measure used to evaluate and compare SR algorithms [61]. However, the ability of MSE (and PSNR) to capture perceptually relevant differences, such as high texture detail, is very limited as they are defined based on pixel-wise image differences [60, 58, 20]. This is illustrated in Figure 2, where highest PSNR does not necessarily reflect the perceptually better SR result. The

Steps of SRGAN

1. A high-resolution image (HR) is downsampled to a low-resolution image (LR)
2. A GAN generator upsamples LR images to super-resolution images (SR)
3. We use a discriminator to distinguish the HR images and backpropagate the GAN loss to train the discriminator and the generator



Generator Network



Discriminator Network




Figure 4: Architecture of Generator and Discriminator Network with corresponding kernel size (k), number of feature maps (n) and stride (s) indicated for each convolutional layer.

The Generator Network

```
#Define blocks to build the generator
```

```
def res_block(ip):  
  
    res_model = Conv2D(64, (3,3), padding = "same")(ip)  
    res_model = BatchNormalization(momentum = 0.5)(res_model)  
    res_model = PReLU(shared_axes = [1,2])(res_model)  
  
    res_model = Conv2D(64, (3,3), padding = "same")(res_model)  
    res_model = BatchNormalization(momentum = 0.5)(res_model)  
  
    return add([ip,res_model])
```

```
def upscale_block(ip):
```

```
    up_model = Conv2D(256, (3,3), padding="same")(ip)  
    up_model = UpSampling2D( size = 2 )(up_model)  
    up_model = PReLU(shared_axes=[1,2])(up_model)  
  
    return up_model
```

```
num_res_block=16
```

```
#Generator model
```

```
def create_gen(gen_ip, num_res_block):  
    layers = Conv2D(64, (9,9), padding="same")(gen_ip)  
    layers = PReLU(shared_axes=[1,2])(layers)  
  
    temp = layers  
  
    for i in range(num_res_block):  
        layers = res_block(layers)  
  
    layers = Conv2D(64, (3,3), padding="same")(layers)  
    layers = BatchNormalization(momentum=0.5)(layers)  
    layers = add([layers,temp])  
  
    layers = upscale_block(layers)  
    layers = upscale_block(layers)  
  
    op = Conv2D(3, (9,9), padding="same")(layers)  
  
    return Model(inputs=gen_ip, outputs=op)
```

Discriminator Network

```
def discriminator_block(ip, filters, strides=1, bn=True):

    disc_model = Conv2D(filters, (3,3), strides = strides, padding="same")(ip)

    if bn:
        disc_model = BatchNormalization( momentum=0.8 )(disc_model)

    disc_model = LeakyReLU( alpha=0.2 )(disc_model)

    return disc_model
```

```
#Discriminator, as described in the original paper
def create_disc(disc_ip):

    df = 64

    d1 = discriminator_block(disc_ip, df, bn=False)
    d2 = discriminator_block(d1, df, strides=2)
    d3 = discriminator_block(d2, df*2)
    d4 = discriminator_block(d3, df*2, strides=2)
    d5 = discriminator_block(d4, df*4)
    d6 = discriminator_block(d5, df*4, strides=2)
    d7 = discriminator_block(d6, df*8)
    d8 = discriminator_block(d7, df*8, strides=2)

    d8_5 = Flatten()(d8)
    d9 = Dense(df*16)(d8_5)
    d10 = LeakyReLU(alpha=0.2)(d9)
    validity = Dense(1, activation='sigmoid')(d10)




    return Model(disc_ip, validity)
```




Compile Model

- VGG19 - weights using “imagenet” & selected layers
- low resolution image 32x32
- high resolution image 128x128
- compile using “binary_crossentropy”
- For each epoch:
 - Train Discriminator - two losses are captured after supplying mix of fake lr images and real lr to the generator, and the the two losses are averaged
 - Train Generator - halt the discriminator and get the image features from vgg prediction on hr images to get the generative loss


Model: "model_2"			
Layer (type)	Output Shape	Param #	
input_3 (InputLayer)	[(None, 128, 128, 3)]	0	
block1_conv1 (Conv2D)	(None, 128, 128, 64)	1792	
block1_conv2 (Conv2D)	(None, 128, 128, 64)	36928	
block1_pool (MaxPooling2D)	(None, 64, 64, 64)	0	
block2_conv1 (Conv2D)	(None, 64, 64, 128)	73856	
block2_conv2 (Conv2D)	(None, 64, 64, 128)	147584	
block2_pool (MaxPooling2D)	(None, 32, 32, 128)	0	
block3_conv1 (Conv2D)	(None, 32, 32, 256)	295168	
block3_conv2 (Conv2D)	(None, 32, 32, 256)	590080	
block3_conv3 (Conv2D)	(None, 32, 32, 256)	590080	
block3_conv4 (Conv2D)	(None, 32, 32, 256)	590080	
Total params: 2,325,568			
Trainable params: 2,325,568			
Non-trainable params: 0			
None			
Model: "model_3"			
Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 32, 32, 3)]	0	[]
model (Functional)	(None, 128, 128, 3)	2044291	['input_1[0][0][*]']
input_2 (InputLayer)	[(None, 128, 128, 3)]	0	[]
model_1 (Functional)	(None, 1)	38249281	['model[0][0][*]']
model_2 (Functional)	(None, 32, 32, 256)	2325568	['model[0][0][*]']
Total params: 42,619,140			
Trainable params: 2,040,067			
Non-trainable params: 40,579,073			
100% [████████████████████] 167/167 [12:00:00<00, 4.32s/it]			
epoch: 1 g_loss: 127.99546988163931 d_loss: [0.02951705 0.66167665]			
100% [████████████████████] 167/167 [11:40:00<00, 4.23s/it]			
epoch: 2 g_loss: 86.70631530030798 d_loss: [0.6692774 0.89221557]			
100% [████████████████████] 167/167 [11:56:00<00, 4.29s/it]			
epoch: 3 g_loss: 81.0359460042425 d_loss: [0.77218723 0.88223353]			
100% [████████████████████] 167/167 [11:57:00<00, 4.30s/it]			
epoch: 4 g_loss: 74.73145766942372 d_loss: [2.22943178 0.75149701]			
100% [████████████████████] 167/167 [12:02:00<00, 4.33s/it]			
epoch: 5 g_loss: 71.64772969951723 d_loss: [0.80506992 0.89526946]			
100% [████████████████████] 167/167 [12:14:00<00, 4.40s/it]			
epoch: 6 g_loss: 69.24715004424135 d_loss: [0.86259081 0.84431138]			
100% [████████████████████] 167/167 [11:58:00<00, 4.30s/it]			
epoch: 7 g_loss: 68.03483305148735 d_loss: [0.47137236 0.87724551]			
100% [████████████████████] 167/167 [12:36:00<00, 4.35s/it]			
epoch: 8 g_loss: 65.0554550826816 d_loss: [0.52766384 0.9011976]			
100% [████████████████████] 167/167 [11:38:00<00, 4.18s/it]			
epoch: 9 g_loss: 67.89687893128538 d_loss: [0.96622361 0.87125749]			
100% [████████████████████] 167/167 [11:21:00<00, 4.08s/it]epoch: 10 g_loss: 64.490			
WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to			

Results 1:







LR Image




Superresolution




Orig. HR image



LR Image




Superresolution




Orig. HR image




LR Image




Superresolution




Orig. HR image



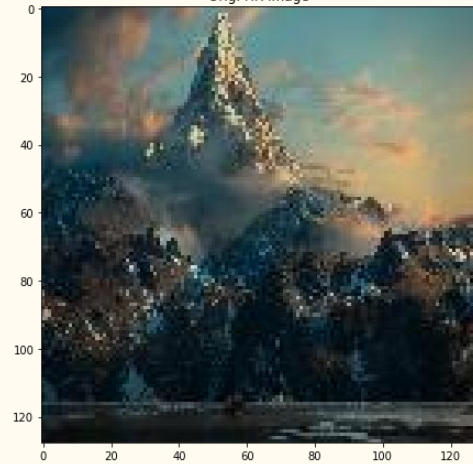
LR Image



Superresolution



Orig. HR image



~The End~

Questions?